

Operating Systems & Concurrency

Week 3 - Threads



Dr. Thiago Braga Rodrigues

tbrodrigues@ait.ie


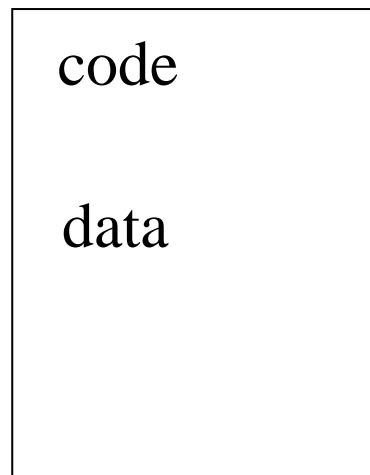
Athlone Institute of Technology
2020

Processes

- In most operating systems, each process has its own private address space and a single thread of control.
- Processes can be switched on and off the CPU by the operating system.
- The OS saves (and restores) the volatile environment (registers), memory management information (location of page table directory) and information about open files for processes.

Multiple Processes

Processes



OS



Process Control Block

Multiple Processes

- Machine code instructions in multiple processes can be arbitrarily interleaved by the runtime environment (the operating system).
- Normally processes don't share memory.

Multiple Threads


- A thread is like a mini-process inside a single (Java) program.
- Machine code instructions in multiple threads can be arbitrarily interleaved by the runtime environment (the JVM or the OS for C++ programs).
- Threads normally do share memory.
 - This has enormous repercussions. Examples?
- How many threads can run at a time in a PC?
 - Hardware and software support!

Multiple Threads


- Threads now have a thread state, i.e.
 - ready, running, blocked
- When one thread inside a process is blocked, another thread can be running.
- [Two threads or processes running at the same time are said to be running concurrently.]

Multiple processes

Process 1




Process 2



Multiple Threads

Process 1




Multiple Threads (cont)

- A thread is a single sequential flow of control through a program.
- One thread can execute for a while.
- Then there is a thread switch (mini-context switch), and its volatile environment is saved.
- Another thread executes.
- Operations in multiple threads can be arbitrarily interleaved by the runtime environment (JVM).

Multiple Threads (cont)

- A process may be multithreaded where same program contain multiple concurrent threads.
- All of the process threads share same memory and open files
 - Each thread gets its own stack, instruction pointer, and registers.
- An OS keep track of threads and stored its per thread information in a data structure called thread control block (TCB).



Example - Multi-threaded File Server

- Two HTTP requests for file are handled simultaneously by two threads.
- First Thread
 - reads request from client
 - issues I/O request for requested file
 - goes to sleep (each Thread has its own Thread state)
- Second Thread can now run
 - reads request from client
 - issues I/O request for requested file
 - goes to sleep
- So multiple requests can be handled at the same time.
- Each thread has its own state. One can be blocked while the other is running

Advantages of threads



Advantages of using threads

- Threads are more efficient
 - The OS does not need to create a new memory map for a new thread (like a process).
 - Does not need to keep track of the state of open files.
- Makes certain programming easy
 - Shared memory makes it trivial to share data among threads
 - Global and static variables can be read and written among all threads in a process
- A multithreaded application can scale performance as the number of processors or cores increases in a system
 - More threads

Threads in Java



Threads in Java

- Java supports the existence of multiple threads in the one program.
- There are two ways to provide a body for a thread.
- The first is to subclass the Thread class and override the Thread class's run() method.
- [The second is to override the Runnable interface.]

Example 1

```
public class SimpleThread extends Thread {  
    public SimpleThread(String str) {  
        super(str);  
    }  
    public void run() {  
        for (int i = 0; i < 10; i++) {  
            System.out.println(i + " " + getName());  
            try {  
                sleep((long)(Math.random() * 1000));  
            } catch (InterruptedException e) {}  
        }  
        System.out.println("DONE! " + getName());  
    }  
}
```


Starting Two SimpleThreads

```
class TwoThreads1 {  
    public static void main (String[] args) {  
        Thread t1 = new SimpleThread1("Jamaica") ;  
        Thread t2 = new SimpleThread1("Fiji") ;  
        t1.start();  
        t2.start();  
    }  
}
```

SimpleThread1

- The class SimpleThread1 contains
 - a constructor for the class which calls the constructor from the super class.
 - a run() method.

```
public Thread(String name)
// creates a new thread with a specified name
public void run()
// subclasses of Thread should override this method.
```

Multiple Threads

- The above example gives two independent threads running at the same time.
- It is completely up to the Java interpreter (JVM) when each thread runs.
- The JVM can time-slice the two threads, like the OS does with processes.
- Up the size of the loop and see what happens.

Multiple Threads

- You can not and should not make any assumptions about when each thread will run.

Thread Methods

- `start()` - makes the thread runnable.
- `run()` - called by JVM when it is the turn of the thread to run.

Example 2

```
public void run() {  
    for (int i = 0; i < 10; i++) {  
        System.out.println(i + " " + getName());  
        try {  
            sleep(1000);  
        } catch (InterruptedException e) {}  
    }  
    System.out.println("DONE! " + getName());  
}
```

sleep(), yield()

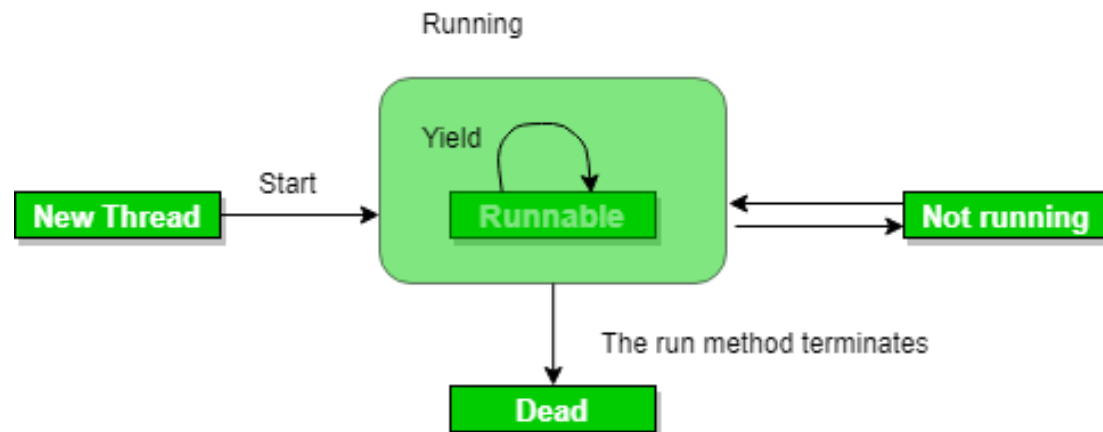
- sleep() – causes the currently executing thread to pause for a specified number of milliseconds.
- When one thread sleeps for a second, the other thread is scheduled.
- yield() – causes the thread to temporarily pause and allow other threads to run.

Thread Methods

- `void run()`
 - body of the thread implemented by programmer
- `void start()`
 - make the thread runnable (called by programmer)
- `void sleep(int milliseconds)`
 - block the thread for a number of milliseconds

Thread Methods

- `void yield()`
 - temporarily stop and let any other runnable thread run
- `join()`
 - `t.join()` - stop the current (parent) thread until thread `t` finishes



Creating Threads - Implementing the Runnable Interface



The Runnable Interface

- The second way to create a Thread in Java is to implement the Runnable Interface.
 - An interface that is to be implemented by a class whose instances are intended to be executed by a thread
- Write a class that implements the interface Runnable.
- It should implement the method `public void run()`.

The Runnable Interface

- To start a Thread, construct a Thread object using the constructor
 - `Thread(Runnable target, String name)`
- Then invoke `start()` on this Thread.

Example (cont)

```
class TwoThreadsRunnable {  
    public static void main (String args[]) {  
        Runnable s1 = new SimpleThreadRunnable("Jamaica") ;  
        Thread t1 = new Thread(s1) ;  
        Runnable s2 = new SimpleThreadRunnable("Fiji") ;  
        Thread t2 = new Thread(s2) ;  
        t1.start();  
        t2.start();  
    }  
}
```

Example

```
class SimpleThreadRunnable implements Runnable {  
    String name ;  
  
    public SimpleThreadRunnable(String name) {  
        this.name = name ;  
    }  
  
    // implement run() as before.
```

Threads - Summary

A decorative graphic consisting of several horizontal stripes in shades of blue and white, located at the bottom of the slide.

Summary

- It is possible to write a program that consists of a number of threads.
- Threads are like mini-processes.
- The main difference between process and a thread is that threads normally share data (variables) while processes normally don't.
- When a number of Java threads execute, the JVM can decide when and for how long to run each thread.

Summary (cont)

- How instructions from different threads are interleaved can not be determined and will normally vary from one execution of the program to the next.
- To create a thread in Java write a class that
 - extends the class Thread
 - implements the method public void run()

Summary (cont)

- Some thread methods
 - `start()` – make the thread runnable
 - `run()` – implement this method
 - `sleep(long time)` – block the thread
 - `yield()` – let another thread run
 - `join()` – wait for a thread to finish

Summary (cont)

- A second way to create a Thread in Java is to
 - Write a class that implements the interface Runnable
 - [implement the run() method]
 - Construct a thread object and pass a Runnable object to the Thread constructor.



Lab